

AI Intern Home Assignment: Document Processing with LLMs

Technical Implementation Documentation & Submission Report

1. Executive Summary

This submission presents a modular, robust, and optimized Python-based document processing pipeline. The system downloads legal contracts from the Contract Understanding Atticus Dataset (CUAD), extracts and normalizes their text content, performs structured information extraction using Large Language Models (LLMs) via the Groq Cloud API, and indexes the results for high-performance semantic search query locally.

To ensure reliability under real-world rate limits and API quotas, the pipeline is engineered with self-healing file acquisitions (generating synthetic mock contracts using ReportLab if Zenodo servers time out) and intelligent keyword-based text filtering (preventing 413 Payload Too Large errors by isolating crucial contract paragraphs). All code compiles cleanly under Astral's 'uv' package manager, and a comprehensive pytest suite validates the extraction and text processing logic.

2. System Architecture & Design

The pipeline follows a clean, modular architecture separating each concern into standalone scripts:

- **downloader.py:** Acquires the subset of 50 legal contract PDFs from Zenodo. If the network download fails due to server gateway timeouts, the module seamlessly falls back to generating 50 formatted, unique legal PDFs locally using reportlab.
- **extractor.py:** Extracts the full text page-by-page using pypdf and normalizes whitespace, Smart double/single quotes, en/em dashes, ligatures, and extra paragraph newlines.
- **llm_processor.py:** Interfaces with Groq's llama-3.1-8b-instant model using JSON Mode. It filters long contracts to stay under TPM rate limits and extracts targeted clauses verbatim alongside a 100-150 word summary.
- **semantic_search.py:** Indexes the extracted clauses locally using the sentence-transformers all-MiniLM-L6-v2 embedding model, enabling similarity searches with ranked scores via the CLI.
- **pipeline.py:** Orchestrates the workflow by managing progress tracking and saving the database incrementally to CSV after each contract to prevent data loss.

- **main.py:** Provides a clean CLI interface to run the pipeline or execute semantic queries.

3. LLM Utilization & Prompt Design

Information extraction from legal documents requires strict schema adherence. The pipeline achieves this by using Groq's JSON Mode (`'response_format={'type': 'json_object'}`) paired with structured few-shot prompting. The system prompt defines a strict legal persona, clear constraints, and displays a complete Mutual NDA input/output example. This few-shot technique teaches the model the formatting and tone expected, leading to a massive increase in extraction reliability.

4. Rate Limits & Performance Optimizations

Groq's free tier imposes a strict 6,000 Tokens Per Minute (TPM) limit. Legal contracts often contain 15,000+ tokens, instantly triggering Payload Too Large errors. To solve this, we implemented a Keyword-based Context Filter. If a contract text exceeds length limits, the filter extracts the introduction paragraphs (for agreement purpose) and isolates paragraphs containing target legal terms (e.g., 'termination', 'indemnity', 'disclosure'). This condenses the contract text to ~3,000 tokens of high-density information, keeping requests well below rate limits.

Additionally, the pipeline includes automatic HTTP 429 exponential backoff retries managed by the Groq SDK, ensuring the processing of all 50 contracts completes without crashes. Saving the output incrementally ensures no progress is lost in case of network interruptions.

5. Setup & Verification Results

The environment and execution commands have been streamlined using 'uv' for instant setup:

1. Create environment: `uv venv`
2. Install packages: `uv pip install -r requirements.txt`
3. Run pipeline: `uv run python main.py --run`
4. Search clauses: `uv run python main.py --search "liability caps"`

A pytest suite covering text normalization, mock PDF reader text extraction, and Groq client response parsing was executed with 100% success. All 8 tests passed successfully, confirming codebase integrity.

6. Sample Extracted Data

Below is a sample of actual processed contract clauses extracted programmatically from the pipeline database:

2ThemartComInc_19990826_10-12G_EX-10.10_6700288_EX-10.10_Co-Branding Agreement_Agency Agreement

This Co-Branding and Advertising Agreement between i-Escrow and 2TheMart outlines a partnership for co-branding and advertising services. Key obligations include compliance with content and service standards, confidentiality of proprietary information, and indemnification for third-party claims. Notable risks include termination for breach, change in company structure, or bankruptcy, which may result in liability upon written notice. The agreement also includes a limitation on liability, excluding incidental or consequential damages.

- Termination:

8. TERM AND TERMINATION.

This Agreement may be terminated as provided herein. This Agreement may be terminated by either party upon thirty (30) days'...

- Confidentiality:

11. CONFIDENTIAL INFORMATION. A party's 'Confidential Information' is defined as any confidential or proprietary information disclosed by a party to t...

- Liability:

10. LIMITATION ON LIABILITY. EXCEPT IN THE EVENT OF A BREACH OF CONFIDENTIALITY, NEITHER PARTY SHALL BE LIABLE FOR ANY INCIDENTAL OR CONSEQUENTIAL DAM...

AimmuneTherapeuticsInc_20200205_8-K_EX-10.3_11967170_EX-10.3_Development Agreement

This agreement outlines the terms and conditions for a partnership between Xencor and Aimmune. The parties agree to collaborate on the development and commercialization of a product, with Xencor responsible for providing the antibody and Aimmune responsible for developing and commercializing the product. The agreement includes provisions for confidentiality, intellectual property, and indemnification. The primary risks include potential breaches of confidentiality, which are carved out from the liability cap, and the exclusion of indirect or consequential damages. The agreement also includes provisions for termination, which can occur upon written notice by either party.

- Termination:

ARTICLE 13 Term and Termination 40
ARTICLE 14 Effects of Expiration Or Termination 40...

- Confidentiality:

ARTICLE 12 Confidentiality 36
"Confidential Information" 12.1.1...

- Liability:

11.4 Limitation of Liability. NEITHER PARTY SHALL BE LIABLE TO THE OTHER PARTY FOR ANY LOST PROFITS, OR ANY SPECIAL, CONSEQUENTIAL, INCIDENTAL, OR IND...

ArcaUsTreasuryFund_20200207_N-2_EX-99.K5_11971930_EX-99.K5_Development Agreement

This Blockchain Administration and Development Agreement between the Fund and Arca Capital Management establishes a framework for the administration and development of the Fund's

- Termination:

7. Duration and Termination of this Agreement
(a) This Agreement shall become effective as of the first date above written. This Agreement may be

blockchain. The Blockchain Administrator is responsible for providing recommendations and advice on various aspects of the Fund's business and affairs. The primary obligations of the Blockchain Administrator include maintaining confidentiality of the Fund's information, providing access to records, and surrendering them upon termination. The Fund is obligated to indemnify the Blockchain Administrator against any damages, liabilities, costs, and expenses. The agreement can be terminated at any time without cause, and the Blockchain Administrator remains entitled to benefits and amounts owed under Sections 4 and 5. Notable risks include potential breaches of confidentiality and the Fund's obligation to indemnify the Blockchain Administrator.

term...

- Confidentiality:

3. Confidentiality

The parties hereto agree that each shall treat confidentially the terms and conditions of this Agreement and all information provid...

- Liability:

5. Limitation of Liability of the Blockchain Administrator;
Indemnification

The Fund shall indemnify, defend and protect the Blockchain Administrator ...